

Tidy Survival Analysis: Applying R's Tidyverse to Survival Data

Module 1. Introduction

LU MAO

hmao@biostat.wisc.edu

Department of Biostatistics & Medical Informatics

University of Wisconsin-Madison

Aug 3, 2025

Table of contents

- Basics of Survival Analysis
- German Breast Cancer Study: A Working Example
- Standard Analysis with `survival` Package
- Summary

Basics of Survival Analysis

Time-to-Event Data


- A common type of outcome in medical and clinical studies
 - **Starting point:** Randomization, diagnosis, enrollment, birth, etc.
 - **Endpoint (Event of interest):** Death, disease onset, hospitalization, etc.
 - **Engineering:** Failure times of machines or components (reliability)
 - **Social sciences:** Time to job change, dropout, or event occurrence
- **Right censoring:**
 - Event not observed within the follow-up period
 - Due to study ending, dropout, or loss to follow-up
 - We only know:

$$T > C$$

where T is event time and C is censoring time

Follow-up (Swimmer) Plot

- A rat tumorigenicity study



Basic Estimands

- **Survival function:** $S(t) = \Pr(T > t)$
 - Probability subject survives beyond time t

- **Hazard function:**

$$\lambda(t) = \lim_{\Delta t \rightarrow 0} \frac{\Pr(t \leq T < t + \Delta t \mid T \geq t)}{\Delta t}$$

- Instantaneous risk of failure at time t

- **Relationship**

$$S(t) = \exp\left(-\int_0^t \lambda(u) du\right)$$

- **Cumulative hazard function:** $\Lambda(t) = \int_0^t \lambda(u) du$

Observed (Censored) Data

- **Notation:** (X, δ)

- $X = \min(T, C)$: observation time (event or censoring)
- $\delta = I(T \leq C)$: event indicator (1 for event, 0 for censoring)

- **Data format**

```
1 # time = X, status = delta (tidy format)
2     id     time status
3 1     1     5     1
4 2     2     3     0
5 3     3     8     1
6 4     4     2     0
7 5     5     6     1
8 # Alternatively
9     id     time
10 1     1     5
11 2     2     3+
12 3     3     8
13 4     4     2+
14 5     5     6
```

German Breast Cancer Study: A Working Example

German Breast Cancer (GBC) Study

- **Study Information**

- **Population:** 686 patients with node-positive breast cancer
- **Objective:** Assess if tamoxifen + chemo reduces mortality/relapse
- **Baseline info:** Age, tumor size, hormone levels, menopausal status, etc.
- **Follow-up:** Median 44 months
 - 171 deaths → exact times known
 - 515 censored → survival time > censoring time

- **Data sets:**

- Mortality data: https://lmaowisc.github.io/tidysurv/data/gbc_mort.txt
- Mortality + relapse: <https://lmaowisc.github.io/tidysurv/data/gbc.txt>
- Download and save in a `data` folder under your root directory

Data Format (I)

- Death only

```
1 # Load mortality data
2 gbc_mort <- read.table("data/gbc_mort.txt", header = TRUE)
3 # Check the first few rows of the data frame
4 head(gbc_mort)
```

	id	time	status	hormone	age	meno	size	grade	nodes	prog	estrg
1	1	74.819672	0	1	38	1	18	3	5	141	105
2	2	65.770492	0	1	52	1	20	1	1	78	14
3	3	47.737705	1	1	47	1	30	2	1	422	89
4	4	4.852459	0	1	40	1	24	1	3	25	11
5	5	61.081967	0	2	64	2	19	2	1	19	9
6	6	63.377049	0	2	49	2	56	1	3	356	64

```
1 # The data frame 'gbc_mort' contains:
2 # time: time (months) to death or censoring
3 # status: event indicator (1 = death, 0 = censoring)
4 # hormone: Hormone therapy (1 = no, 2 = yes); age: Age at diagnosis (years);
5 # meno: Menopausal status (1 = no, 2 = yes); size: Tumor size (mm); grade: Tumor grade (1-3);
6 # nodes: Number of positive lymph nodes; prog: Progesterone receptor level (fmol/mg); estrg:
7 # Estrogen receptor level (fmol/mg).
```

Data Format (II)

- Mortality + relapse

```
1 # Load mortality + relapse data
2 gbc <- read.table("data/gbc.txt", header = TRUE)
3 # Check the first few rows of the data frame
4 head(gbc)
```

	id	time	status	hormone	age	meno	size	grade	nodes	prog	estrg
1	1	43.83607	1	1	38	1	18	3	5	141	105
2	1	74.81967	0	1	38	1	18	3	5	141	105
3	2	46.55738	1	1	52	1	20	1	1	78	14
4	2	65.77049	0	1	52	1	20	1	1	78	14
5	3	41.93443	1	1	47	1	30	2	1	422	89
6	3	47.73770	2	1	47	1	30	2	1	422	89

```
1 # The data frame 'gbc' contains:
2 # time: time (months) to death, relapse, or censoring
3 # status: event indicator (1 = relapse, 2 = death, 0 = censoring)
4 # other covariates the same as in gbc_mor.
```

Analysis Goals

- **Descriptive**

- Summarize patient characteristics
- Visualize survival distributions

- **Inferential**

- Compare survival curves (e.g., hormone therapy vs. no hormone therapy)
- Assess impact of covariates on survival (e.g., age, tumor size, etc.)
- Model competing risks (e.g., relapse vs. death)

- **Predictive**

- Develop risk prediction models
- Evaluate model performance (e.g., concordance index, calibration)

Standard Analysis with **survival** Package

Survival Package Overview

- Key Functions

- `Surv()`: Create survival object
- `survfit()`: Fit Kaplan-Meier survival curves
- `survdiff()`: Compare survival curves (log-rank test)
- `coxph()`: Fit Cox proportional hazards regression models
- `survreg()`: Fit parametric survival regression models

Kaplan-Meier Survival Curves

- Create dataset for relapse-free survival

```
1 # Sort by subject id, then time
2 o <- order(gbc$id, gbc$time)
3 gbc <- gbc[o,]
4 # Keep only first row per subject => first event
5 df <- gbc[!duplicated(gbc$id), ]
6 # Convert status > 0 to 1 if it is either relapse or death
7 df$status <- ifelse(df$status > 0, 1, 0)
8 head(df)
```

	id	time	status	hormone	age	meno	size	grade	nodes	prog	estrg
1	1	43.836066	1	1	38	1	18	3	5	141	105
3	2	46.557377	1	1	52	1	20	1	1	78	14
5	3	41.934426	1	1	47	1	30	2	1	422	89
7	4	4.852459	0	1	40	1	24	1	3	25	11
8	5	61.081967	0	2	64	2	19	2	1	19	9
9	6	63.377049	0	2	49	2	56	1	3	356	64

Kaplan-Meier Curves (I)

- Fit Kaplan-Meier survival curves

```
1 library(survival)
2 # Fit KM curves by hormone treatment group
3 km_fit <- survfit(Surv(time, status) ~ hormone, data = df)
4 km_fit
```

Call: survfit(formula = Surv(time, status) ~ hormone, data = df)

	n	events	median	0.95LCL	0.95UCL
hormone=1	440	205	50.1	42.5	59.5
hormone=2	246	94	66.2	62.9	NA

Kaplan-Meier Curves (II)

- Summarize survival estimates at specified time points
 - For example, at 6, 12, 24, and 36 months

```
1 summary(km_fit, times = c(6, 12, 24, 36))
```

Call: `survfit(formula = Surv(time, status) ~ hormone, data = df)`


hormone=1							
time	n.risk	n.event	survival	std.err	lower 95% CI	upper 95% CI	
6	419	9	0.979	0.00691	0.966	0.993	
12	379	35	0.897	0.01476	0.868	0.926	
24	280	73	0.720	0.02203	0.678	0.764	
36	195	41	0.606	0.02475	0.559	0.656	

hormone=2							
time	n.risk	n.event	survival	std.err	lower 95% CI	upper 95% CI	
6	236	4	0.983	0.00826	0.967	1.000	
12	223	8	0.950	0.01418	0.922	0.978	
24	177	38	0.785	0.02701	0.733	0.839	
36	136	16	0.708	0.03047	0.650	0.770	

Kaplan-Meier Curves (III)

- Plot Kaplan-Meier survival curves by group

```
1 plot(km_fit, ylim = c(0,1), xlab = "Time (months)", ylab = "Relapse-free survival probability",  
2       col = c("red", "blue"), conf.int = TRUE)  
3 # Add legend  
4 legend(1, 0.2, col=c("red", "blue"), lty = 1,  
5       c("No hormone", "Hormone")) # Legend text
```



Log-Rank Test

- Compare survival curves between groups

```
1 lgr_obj <- survdiff(Surv(time, status) ~ hormone, data = df)
2 lgr_obj # Print log-rank test results
```

Call:

```
survdiff(formula = Surv(time, status) ~ hormone, data = df)
```

	N	Observed	Expected	(O-E) ² /E	(O-E) ² /V
hormone=1	440	205	180	3.37	8.56
hormone=2	246	94	119	5.12	8.56

Chisq= 8.6 on 1 degrees of freedom, p= 0.003

```
1 lgr_obj$pvalue # Extract p-value
```

```
[1] 0.003427282
```

Exercise

Perform a log-rank test on treatment stratified by patient menopausal status `meno`.

► Solution

Cox Model - Model Specification

- **Cox proportional hazards model**

$$\lambda(t \mid Z) = \lambda_0(t) \exp(\beta_1 Z_1 + \beta_2 Z_2 + \dots + \beta_p Z_p)$$

- $\lambda_0(t)$: baseline hazard function
- $Z = (Z_1, \dots, Z_p)^T$: covariates (e.g., hormone therapy, age, tumor size)
- $\beta = (\beta_1, \dots, \beta_p)^T$: regression coefficients
- $\exp(\beta_j)$: hazard ratio for covariate Z_j

- **Proportional hazards (PH) assumption**

$$\frac{\lambda(t \mid Z)}{\lambda(t \mid Z^*)} = \exp\{\beta^T (Z - Z^*)\}$$

- HR constant over time, i.e., $\beta(t) \equiv \beta$ (for each covariate)

Cox Model - Model Fitting (I)

- Model fitting: `survival::coxph()`

```
1 cox_fit <- coxph(Surv(time, status) ~ hormone + meno + age + grade + size + prog + estrg,  
2                   data = df)  
3 summary(cox_fit) # Print model summary
```

Call:

```
coxph(formula = Surv(time, status) ~ hormone + meno + age + grade +  
      size + prog + estrg, data = df)
```

n= 686, number of events= 299

	coef	exp(coef)	se(coef)	z	Pr(> z)	
hormone	-0.3422139	0.7101963	0.1290669	-2.651	0.00801	**
meno	0.2765637	1.3185909	0.1837781	1.505	0.13236	
age	-0.0087813	0.9912572	0.0093375	-0.940	0.34700	
grade	0.2785797	1.3212519	0.1051531	2.649	0.00807	**
size	0.0152793	1.0153966	0.0036877	4.143	3.42e-05	***
prog	-0.0023307	0.9976720	0.0005803	-4.016	5.91e-05	***
estrng	0.0001678	1.0001679	0.0004669	0.359	0.71923	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
exp(coef) exp(-coef) lower 95 upper 95
```

Cox Model - Model Fitting (II)

- Extracting $\hat{\beta}$ and $\text{var}(\hat{\beta})$

```
1 beta <- cox_fit$coefficients # Estimated coefficients
2 vbeta <- vcov(cox_fit) # Estimated variance-covariance matrix
3 # Extract regression table (as data frame)
4 coef(summary(cox_fit))
```

	coef	exp(coef)	se(coef)	z	Pr(> z)
hormone	-0.3422138547	0.7101963	0.1290669350	-2.6514448	8.014821e-03
meno	0.2765636858	1.3185909	0.1837780595	1.5048787	1.323553e-01
age	-0.0087812621	0.9912572	0.0093375120	-0.9404285	3.469978e-01
grade	0.2785796730	1.3212519	0.1051531448	2.6492757	8.066449e-03
size	0.0152793172	1.0153966	0.0036877471	4.1432660	3.423945e-05
prog	-0.0023307288	0.9976720	0.0005803186	-4.0162919	5.912102e-05
estrg	0.0001678465	1.0001679	0.0004669057	0.3594870	7.192308e-01

- **Conclusion**

- Hormone therapy significantly reduces the risk of relapse or death by $1 - 0.710 = 29\%$ ($p = 0.008$)

Cox Model - Prediction (I)

- Predicted survival function

$$\hat{S}(t | z) = \exp\left\{-\exp(\hat{\beta}^T z)\hat{\Lambda}_0(t)\right\}$$

- Prepare new data for prediction

```
1 # Create new data for prediction
2 # specify all covariate values
3 new_data <- data.frame(hormone = 1, meno = 1,
4                         age = 45, grade = 2,
5                         size = 20, prog = 100,
6                         estrg = 100)
7 new_data
```

	hormone	meno	age	grade	size	prog	estr
1	1	1	45	2	20	100	100

Cox Model - Prediction (II)

- Predict survival probabilities at specified time points

```
1 # Predict survival probabilities at 6, 12, 24, 26 months
2 predicted_survival <- survfit(cox_fit, newdata = new_data[1, ], times = c(6, 12, 24, 36))
3 summary(predicted_survival, times = c(6, 12, 24, 36))
```


Call: survfit(formula = cox_fit, newdata = new_data[1,], times = c(6, 12, 24, 36))

time	n.risk	n.event	survival	std.err	lower 95% CI	upper 95% CI
6	655	13	0.985	0.00441	0.976	0.994
12	602	43	0.933	0.01059	0.913	0.954
24	457	111	0.786	0.02304	0.743	0.833
36	331	57	0.696	0.02925	0.641	0.755

Cox Model - Prediction (III)

- Plot predicted survival function

```
1 # Plot predicted survival function
2 plot(predicted_survival, ylim = c(0, 1), xlab = "Time (months)",
3       ylab = "Relapse-free survival probability", conf.int = TRUE)
```



Cox Model - Check PH Assumptions (I)

- Schoenfeld residuals

- Difference between observed and expected covariate values at each event time
- Use `cox.zph()` to test PH assumption

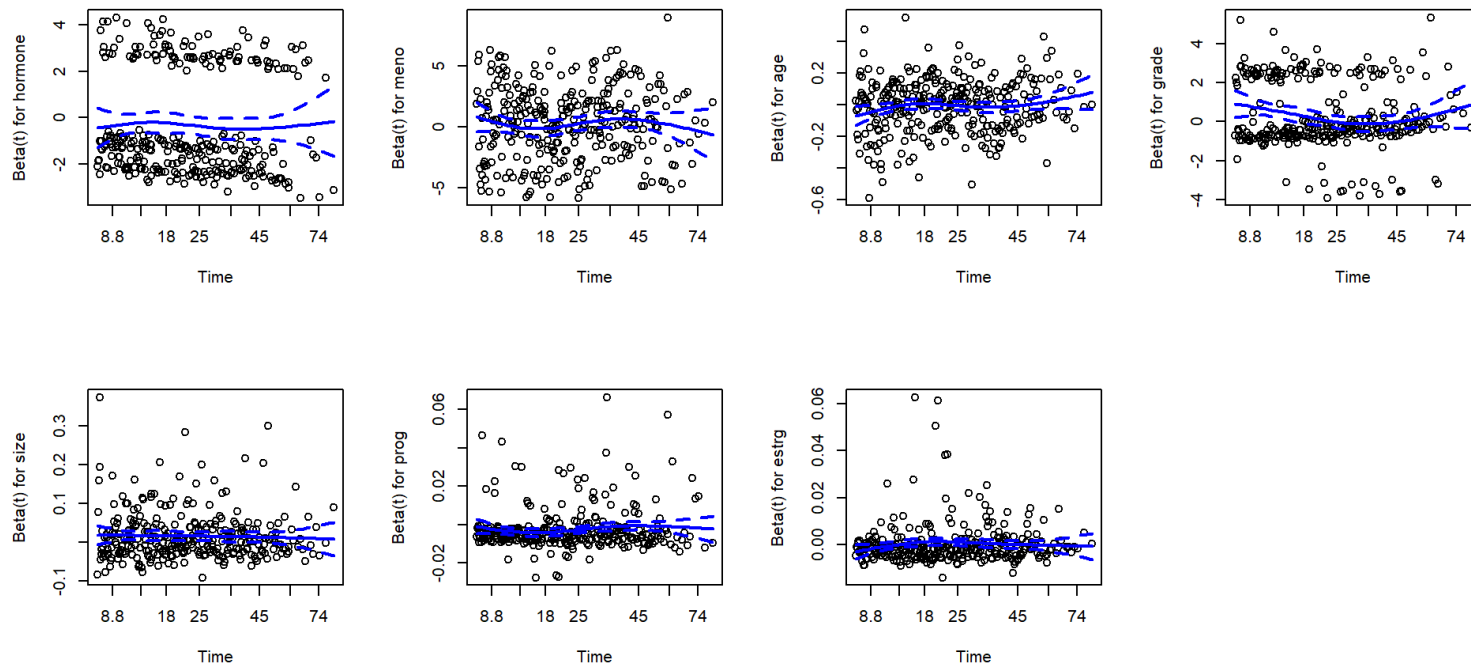
```
1 ph_test <- cox.zph(cox_fit)
2 ph_test # Print test results
```

	chisq	df	p
hormone	0.272	1	0.6017
meno	5.514	1	0.0189
age	9.430	1	0.0021
grade	8.490	1	0.0036
size	0.872	1	0.3505
prog	4.881	1	0.0272
estrg	5.403	1	0.0201
GLOBAL	20.636	7	0.0043

Cox Model - Check PH Assumptions (II)

- Graphical check of PH assumptions
 - Plot Schoenfeld residuals against time

```
1 par(mfrow= c(2, 4)) # Set up 2x4 plotting area for 7 covariates
2 plot(ph_test, se = TRUE, col = "blue", lwd = 2) # Plot Schoenfeld residuals
```




Cox Model - Check Covariate Forms

- Check linearity of covariate effects
 - Plot martingale residuals against (quantitative) covariates

```
1 # Extract martingale residuals
2 mart_resid <- residuals(cox_fit, type = 'martingale')
```

► Plotting



Coding Exercise

Exercise

Residual analyses show that the proportional hazards assumption is violated for tumor grade, and that the effect of age is not linear.

Fit a different model to address these issues.

► Sample solution

Summary

Key Takeaways

- Survival analysis is essential for (often censored) time-to-event data
- Key estimands: survival function, hazard function, cumulative hazard
- Standard analysis tools
 - Kaplan-Meier curves (`survfit()`)
 - Log-rank test (`survdifff()`)
 - Cox proportional hazards model (`coxph()`)

Open Questions

- Efficient/effective presentation of survival probabilities
 - Point estimates, confidence intervals
- Customizable survival curves
 - Add at risk table below graph
- Presentation of regression results
 - Hazard ratios, confidence intervals, p-values
 - Visualize regression results (e.g., forest plots)