

Tidy Survival Analysis: Applying R's Tidyverse to Survival Data

Module 2. Data Manipulation with Tidyverse

LU MAO

hmao@biostat.wisc.edu

Department of Biostatistics & Medical Informatics

University of Wisconsin-Madison

Aug 3, 2025

Table of contents

- Overview of Tidyverse
- Tidying Survival Data
- Visualizing Subject Follow-Up
- Creating “Table 1”
- Summary

Overview of Tidyverse

The **tidyverse** Ecosystem

- **Motivation:** tidy data for reproducible analysis
- **Key packages**
 - **dplyr** (filtering, mutating, grouping, summarizing)
 - **tidyr** (pivoting, nesting, reshaping)
 - **tibble** (modern data frames)
 - **readr** / **haven** (importing .csv or .sas7bdat)
 - **lubridate** (handling time variables)
 - **ggplot2** (visualization)

```
1 # Load core tidyverse packages
2 library(tidyverse)
```

Basic Functionalities

- **Data manipulation:** using `dplyr` verbs
 - `mutate()` to create new variables (e.g., age group, log-transformed labs)
 - `filter()` to subset by treatment or age
 - `select()` and `rename()` for variable formatting
 - `arrange()` to sort
 - `group_by()` and `summarize()` for descriptive summaries by arm
- **Data reshaping:** using `tidyr` functions
 - `pivot_longer()` to convert wide to long format
 - `pivot_wider()` to convert long to wide format
 - `nest()` and `unnest()` for hierarchical data

A Simple Example

- Example dataset

```
1 # Simulated data example
2 df1 <- tibble(
3   id = 1:6,
4   trt = c("A", "A", "B", "B", "A", "B"),
5   age = c(65, 70, 58, 60, 64, 59),
6   time = c(5, 8, 12, 3, 2, 6),
7   status = c(1, 0, 1, 1, 0, 0) # 1 = event, 0 = censored
8 )
9 df1
```

A tibble: 6 × 5

	id	trt	age	time	status
	<int>	<chr>	<dbl>	<dbl>	<dbl>
1	1	A	65	5	1
2	2	A	70	8	0
3	3	B	58	12	1
4	4	B	60	3	1
5	5	A	64	2	0
6	6	B	59	6	0

Native Pipe Operator: |>

- What is |>

- Introduced in R 4.1 (hot key: `Ctrl + Shift + M`)
- Passes the result of one expression into the first argument of the next
- Same idea as `%>%`, but **built into base R**

- Example

```
1 df1 |> # passes tibble data frame df1 to the next function
2   mutate(age_group = if_else(age >= 65, "older", "younger")) |> # create age group
3   filter(trt == "A") |> # filter for treatment A
4   arrange(time) # sort by time
```

A tibble: 3 × 6

	id	trt	age	time	status	age_group
	<int>	<chr>	<dbl>	<dbl>	<dbl>	<chr>
1	5	A	64	2	0	younger
2	1	A	65	5	1	older
3	2	A	70	8	0	older

Summarizing and Grouping

- Survival-specific summaries (e.g., number of events)
 - `group_by()` and `summarize()` for descriptive summaries by arm

```
1 df1 |>
2   group_by(trt) |> # group by treatment arm
3   summarize( # summarize each group
4     n = n(), # count number of rows (subjects)
5     events = sum(status), # sum of events (status = 1)
6     median_time = median(time) # median survival time
7   )
```

A tibble: 2 × 4

	trt	n	events	median_time
	<chr>	<int>	<dbl>	<dbl>
1	A	3	1	5
2	B	3	2	6

What Does “Tidy” Mean?

A dataset is tidy if:

- Each **variable** is a column
- Each **observation** is a row
- Each **type of observational unit** is a table

— Hadley Wickham, *Tidy Data* (2014)

<https://www.jstatsoft.org/article/view/v059i10>

Why Tidy Data?

- **Tidy data principles**

- Easy to reshape and transform
- Compatible with `ggplot2`, `dplyr`, `tidyr`, and modeling tools
- Encourages modular and reproducible code

- **Messy data challenges:**

- Time in rows, covariates in columns
- Multiple data types in one column
- Separate randomization and event/censoring dates
- Missing/censored values inconsistently coded

Tidy Survival Data

- Possible pre-processing steps

- Calculate survival time from start to event/censoring
- Creating the (X, δ) structure expected by `Surv()`
- Reshaping data to long format in case of multiple events

- An Example

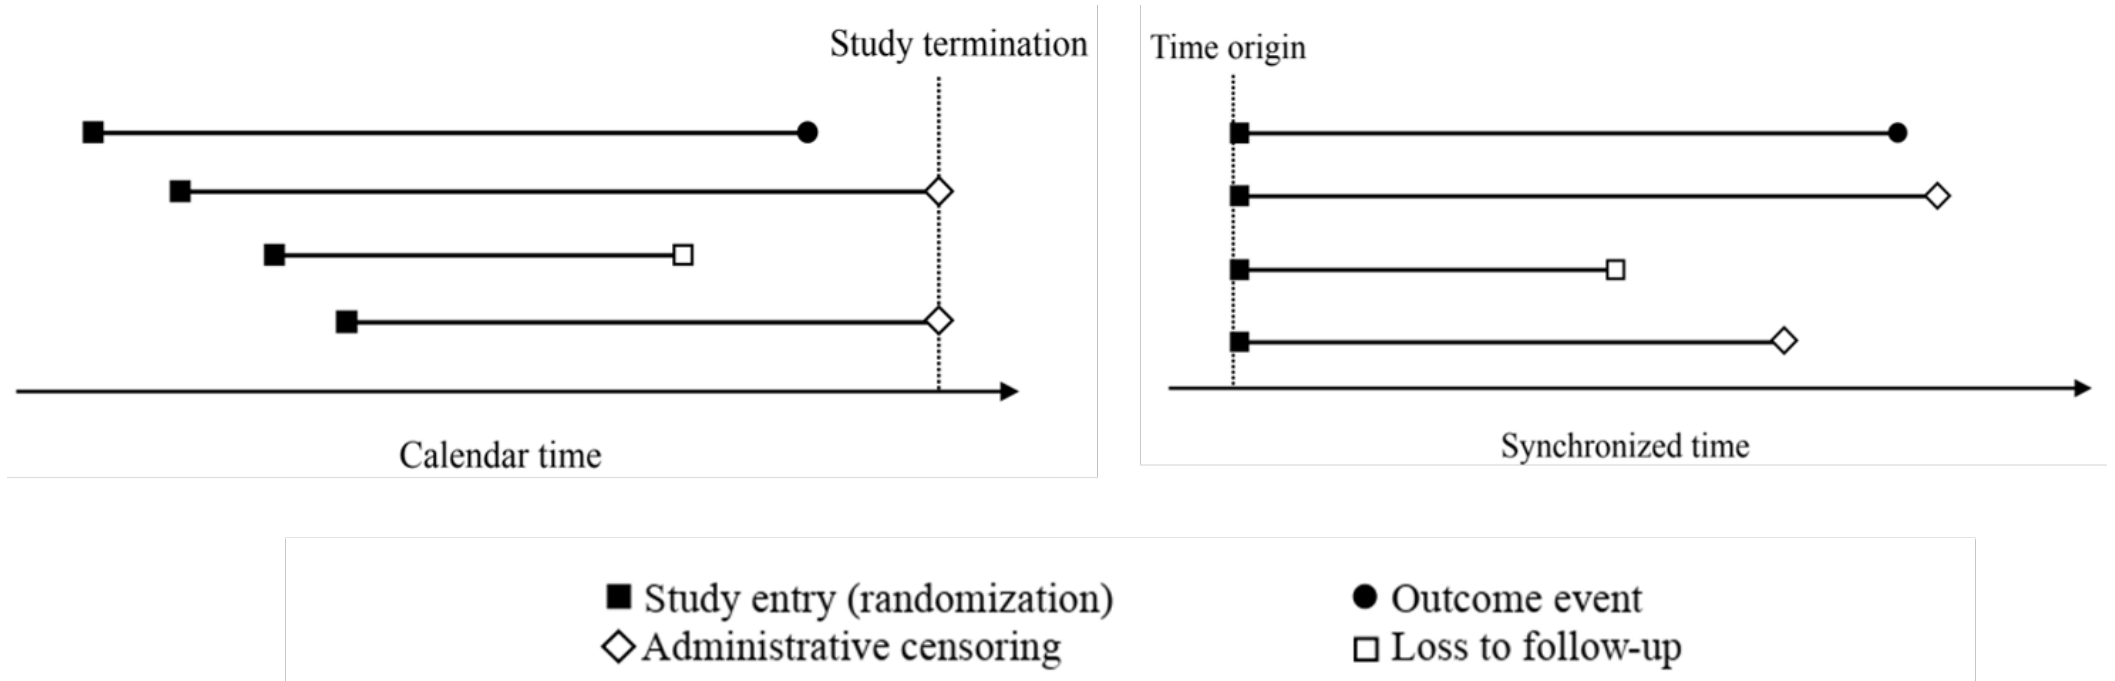
```
1 # Load GBC data
2 gbc <- read.table("data/gbc.txt", header = TRUE)
3 head(gbc)
```

	id	time	status	hormone	age	meno	size	grade	nodes	prog	estrg
1	1	43.83607	1	1	38	1	18	3	5	141	105
2	1	74.81967	0	1	38	1	18	3	5	141	105
3	2	46.55738	1	1	52	1	20	1	1	78	14
4	2	65.77049	0	1	52	1	20	1	1	78	14
5	3	41.93443	1	1	47	1	30	2	1	422	89
6	3	47.73770	2	1	47	1	30	2	1	422	89

Tidying Survival Data

Calendar vs. Event Times

- Time from start to event/censoring ()



Dates to Time Difference

- A data example

```
1 # Example: raw dates as character strings
2 df2 <- tibble(
3   id = 1:3,
4   rand_date = c("2022-01-01", "2022-01-15", "2022-01-20"),
5   end_date = c("2022-04-01", "2022-06-01", "2022-03-15"),
6   status = c("dead", "censored", "dead")
7 )
8 df2
```

A tibble: 3 × 4

	id	rand_date	end_date	status
	<int>	<chr>	<chr>	<chr>
1	1	2022-01-01	2022-04-01	dead
2	2	2022-01-15	2022-06-01	censored
3	3	2022-01-20	2022-03-15	dead

Parsing Dates and Calculating Time

- Using `lubridate` to parse dates

- `ymd()` for “year-month-day” format
- `mdy()` for “month-day-year” format

```
1 # Parse dates and calculate time/status
2 df2 |>
3   mutate(
4     rand_date = ymd(rand_date), # convert character to Date
5     end_date = ymd(end_date), # convert character to Date
6     time = as.numeric(end_date - rand_date), # calculate time in days
7     status = if_else(status == "dead", 1, 0) # convert status to 1/0
8   )
```

A tibble: 3 × 5

	id	rand_date	end_date	status	time
	<int>	<date>	<date>	<dbl>	<dbl>
1	1	2022-01-01	2022-04-01	1	90
2	2	2022-01-15	2022-06-01	0	137
3	3	2022-01-20	2022-03-15	1	54

Exercise: Calculate Survival Time (I)

- Calculate `time` and `status` variables for `df3`:

```
1 # create a df3 with dates in the form of month-day-year
2 df3 <- tibble(
3   id = 1:3,
4   rand_date = c("Jan-01-2022", "01-15-2022", "01-20-2022"),
5   end_date = c("04-01-2022", "Jun-01-2022", "03-15-2022"),
6   status = c("dead", "censored", "dead")
7 )
8 df3
```

A tibble: 3 × 4

	id	rand_date	end_date	status
	<int>	<chr>	<chr>	<chr>
1	1	Jan-01-2022	04-01-2022	dead
2	2	01-15-2022	Jun-01-2022	censored
3	3	01-20-2022	03-15-2022	dead

Exercise: Calculate Survival Time (II)

- Hint: use `mdy()` to parse dates

► Solution

- More about manipulating dates
 - [lubridate](#) official documentation
 - R for Data Science: [Dates and times](#)

Parsing Censored Observations

- Alternative formats for censored times
 - "32+", ">17", etc
 - `parse_number()` for get time; `str_detect()` for status

```
1 # Example data: relapse times with "+" indicating censoring
2 MP <- c(10, "32+", 23, "25+")
3 # Convert to (time, status) format
4 df4 <- tibble(
5   MP = MP,                                # Original data
6   time = parse_number(MP),                # Extract numeric part
7   status = 1 - str_detect(MP, "\\+")      # Censored if "+" detected
8 )
9 df4
```

A tibble: 4 × 3

	MP	time	status
	<chr>	<dbl>	<dbl>
1	10	10	1
2	32+	32	0
3	23	23	1
4	25+	25	0

Exercise: Parse Censored Times

- Task: Parse `MP` in `df5` to create `time` and `status`

```
1 df5 <- tibble(  
2   MP = c(10, "32+", 23, ">25")  
3 )
```

► Solution

- More on string operation
 - [stringr](#) official documentation
 - R for Data Science: [Strings](#)

Reshaping Data

- Why reshape?

- Multiple events per subject
- Wide format (multiple columns) long format (one row per event)

```
1 # Example: wide format with multiple events
2 df6 <- tibble(
3   id = 1:3,
4   prog_time = c(10, 20, 30),
5   prog_status = c(1, 0, 1), # 1 = progression, 0 = censored
6   death_time = c(15, 20, 35),
7   death_status = c(0, 1, 1) # 1 = dead, 0 = censored
8 )
9 # 1: progression at 10, censored at 15
10 # 2: dead at 20 without progression
11 # 3: progression at 30, dead at 35
12 df6
```

A tibble: 3 × 5

	id	prog_time	prog_status	death_time	death_status
	<int>	<dbl>	<dbl>	<dbl>	<dbl>
1	1	10	1	15	0
2	2	20	0	20	1
3	3	30	1	35	1

Wide to Long

- Using `pivot_longer()`
 - Convert wide format to long format
 - Specify `names_to` and `values_to` for new columns

```
1 df7 <- df6 |>
2   pivot_longer(
3     cols = c(prog_time, prog_status, death_time, death_status), # columns to reshape
4     names_to = c("event", ".value"), # .value keeps the variable name, event is the new column
5     names_pattern = "(.*)_(.*)" # split by underscore
6   )
7 df7
```

```
# A tibble: 6 × 4
   id event  time status
<int> <chr> <dbl>   <dbl>
1     1 prog    10       1
2     1 death   15       0
3     2 prog    20       0
4     2 death   20       1
5     3 prog    30       1
6     3 death   35       1
```

Exercise: Clean Up

- **Task:** Clean up `df7` to create a tidy survival dataset
 - Remove rows with `event = prog` and `status = 0` (non-terminal event)
 - Recode `status = 2` for death events


► Solution

- More on reshaping data
 - [tidyr](#) official documentation
 - R for Data Science: [Data tidying](#)

Visualizing Subject Follow-Up

Swimmer Plot

- What is a swimmer plot?
 - Visualizes subject follow-up
 - Each row represents a subject
 - Horizontal lines show time to event/censoring



Swimmer Plot Basics

- Using `ggplot2`
 - `geom_linerange()` for horizontal lines
 - `geom_point()` for events
 - `facet_wrap()` for treatment arms (optional)
- A data example

```
1 # Example data: rat survival times
2 df8 <- tibble(
3   time = c(101, 55, 67, 23, 45, 98, 34, 77, 91, 104, 88),
4   status = c(0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1),
5   group = c("A", "A", "A", "B", "B", "B", "A", "B", "B", "A", "B")
6 ) |>
7   mutate(
8     id = row_number(), # create id column using row number
9     .before = 1 # place id before time
10  )
```

Creating a Swimmer Plot


- Code to reproduce previous plot

```
1 # Specify the plot
2 fig8 <- df8 |>
3   # Set-up: id on the y-axis, time on the x-axis
4   ggplot(aes(x = time, y = reorder(id, time))) + # reorder id by time
5   # Add geometric objects
6   geom_linerange(aes(xmin = 0, xmax = time)) + # horizontal lines from 0 to time
7   # Add points for events/censoring, distinguish by status
8   geom_point(aes(shape = factor(status)), size = 2.5, fill = "white") +
9   # Add vertical line at x = 0
10  geom_vline(xintercept = 0, linewidth = 1) +
11  theme_minimal() + # use minimal theme
12  # Format y axis
13  scale_y_discrete(name = "Rats") + # y-axis label
14  # Format x axis (label, breaks, no expansion on left, 0.05 expansion on right)
15  scale_x_continuous(name = "Time (days)", breaks = seq(0, 100, by = 20),
16                    expand = expansion(c(0, 0.05))) +
17  # Format point shape (pch = 23 for censoring, pch = 19 for event; label shape)
18  scale_shape_manual(values = c(23, 19), labels = c("Censoring", "Tumor development")) +
19  # Further formatting using theme()
20  theme(
21    legend.position = "top", # place legend at the top
22    legend.title = element_blank(), # no legend title
```

Exercise: Swimmer Plot by Group

- **Task:** Create a swimmer plot for `df8` by `group`
 - Use `facet_wrap()` to create separate panels for each group
 - Add a title “Swimmer Plot of Rat Survival Times”

Swimmer Plot of Rat Survival Times



► Solution

Creating “Table 1”

Descriptive Statistics

- Importance of Table 1

- Summarizes baseline characteristics
- Provides context for formal analysis

- Using `gtsummary`

- `tbl_summary()` for descriptive statistics
- `add_p()` for p-values comparing groups (not recommended for randomized trials)
- `add_overall` to add overall summary
- `modify_header()` to customize table headers

Basic Syntax of `tbl_summary()`

- Common arguments

- `by = "group"` to summarize by group
- `include = c("variable1", "variable2")` to include specific variables
- `label = list(variable = "Label")` to customize variable labels
- `statistic = list(variable ~ "statistic")` to specify statistics
 - `statistic = list(all_continuous() ~ "{mean} ({sd})")` for mean and SD
- `digits = list(variable ~ 2)` to set decimal places

A Simple Example

- Example dataset

```
1 # Example data: 10 subjects with treatment, age, and sex
2 df9 <- tibble(
3   id = 1:10,
4   time = c(101, 55, 67, 23, 45, 98, 34, 77, 91, 104),
5   status = c(0, 1, 1, 0, 1, 0, 1, 0, 1, 0), # 0 = censored, 1 = event
6   trt = c("A", "A", "B", "B", "A", "B", "A", "B", "A", "B"),
7   sex = c("M", "F", "M", "F", "M", "F", "M", "F", "M", "F"),
8   age = c(65, 70, 58, 60, 64, 59, 66, 62, 68, 61)
9 )
10 head(df9)
```

A tibble: 6 × 6

	id	time	status	trt	sex	age
	<int>	<dbl>	<dbl>	<chr>	<chr>	<dbl>
1	1	101	0	A	M	65
2	2	55	1	A	F	70
3	3	67	1	B	M	58
4	4	23	0	B	F	60
5	5	45	1	A	M	64
6	6	98	0	B	F	59

Creating a Summary Table

```
1 library(gtsummary) # load package
2 df9 |>
3   tbl_summary(
4     by = trt,          # summarize by treatment arm
5     include = c(sex, age, time, status), # include specific variables
6     label = list(      # label variables
7       time = "Follow-up time (months)",
8       status = "Events"
9     )
10  )
```

Characteristic	A, N = 5 ¹	B, N = 5 ¹
sex		
F	1 (20%)	4 (80%)
M	4 (80%)	1 (20%)
age	66.0 (65.0, 68.0)	60.0 (59.0, 61.0)
Follow-up time (months)	55 (45, 91)	77 (67, 98)
Events	4 (80%)	1 (20%)
¹ n (%); Median (IQR)		

Exercise: Summarize GBC Data (I)

- Task: Summarize the GBC mortality data ([gbc_mort.txt](#)) like below

Characteristic	Hormone, N = 246 ⁷	No Hormone, N = 440 ⁷	Overall, N = 686 ⁷
Follow-up time (months)	48 (29, 61)	41 (25, 57)	44 (26, 60)
Death	56 (23%)	115 (26%)	171 (25%)
Age (years)	58 (50, 63)	50 (45, 59)	53 (46, 61)
Menopausal status	187 (76%)	209 (48%)	396 (58%)
Tumor size (mm)	25 (20, 35)	25 (20, 35)	25 (20, 35)
Tumor grade			
1	33 (13%)	48 (11%)	81 (12%)
2	163 (66%)	281 (64%)	444 (65%)
3	50 (20%)	111 (25%)	161 (23%)
Number of nodes	3 (1, 7)	3 (1, 7)	3 (1, 7)
Progesterone (fmol/mg)	35 (7, 133)	32 (7, 130)	33 (7, 132)
Estrogen (fmol/mg)	46 (9, 183)	32 (8, 92)	36 (8, 114)
⁷ Median (IQR); n (%)			

Exercise: Summarize GBC Data (II)

- **Points to note**

- Summarize by hormone therapy (`hormone`)
- Include variables: `time`, `status`, `age`, `meno`, `size`, `grade`, `nodes`, `prog`, `estrg`
- Label variables appropriately
- Add overall summary column at the end

Exercise: Summarize GBC Data (III)

► Solution

Exercise: Summarize GBC Data (IV)

- **Task:** summarize *relapse* and death data from `gbc.txt`
 - Hint: `group_by(id)` and `summarize()`

Characteristic	Hormone, N = 246 ¹	No Hormone, N = 440 ¹	Overall, N = 686 ¹
Relapse	94 (38%)	205 (47%)	299 (44%)
Death	56 (23%)	115 (26%)	171 (25%)
Composite	94 (38%)	205 (47%)	299 (44%)
Relapse then death	56 (23%)	115 (26%)	171 (25%)
¹ n (%)			

Exercise: Summarize GBC Data (V)

► Solution

Summary

Key Takeaways

- **Tidyverse** provides powerful tools for data manipulation and visualization
- **Tidy data** principles simplify analysis and visualization
- **Survival data** may require pre-processing steps (`dplyr`, `tidyr`, `lubridate`)
- **Swimmer plots** effectively visualize subject follow-up (`ggplot2`)
- **Descriptive statistics** can be easily summarized using `gtsummary::tbl_summary()`

Next Steps

- Format analysis results from the `survival` package:
 - Nonparametric estimates with `survfit()`
 - Regression models with `coxph()`
- Explore advanced visualization techniques:
 - Kaplan–Meier curves with `ggsurvfit` or `survminer`
 - Layered plots using `ggplot2`
 - Annotated plots for publications